

# Nationwide Application Assessment for Computational Telematics

## Part 2

### Jason Barkeloo

Now let's import necessary libraries.

numpy for various linear algebra libraries

pandas for convenient file reading as well as dataframe structures that are convenient to work with

matplotlib for various basic plots

geopy for geodesic distances used for calculations in part 1 sklearn, keras (using locally installed tensorflow) for neural network applications

```
In [135]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from geopy.distance import geodesic
import sys
import sklearn
import keras
from sklearn.model_selection import train_test_split
from keras.models import Model, Sequential
from keras.layers import Dense, Dropout, Input
from keras.callbacks import EarlyStopping
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import confusion_matrix
import xgboost as xgb
```

## Part 2 Modeling

```
In [136]: sim_sum_tot = pd.read_csv(r"C:\Users\JTBar\Documents\Telematics Exercise Files
\simulated_summary_total.csv")
```

In [137]: `sim_sum_tot.head()`

Out[137]:

	Vehicle	Days	Distance	HardBrakes	HardAccelerations	NightTime_Pct	VehicleType	Loss
0	1	365	13114	152	56	0.005	SUV	0
1	2	365	18707	147	1	0.010	SUV	0
2	3	365	16659	151	127	0.019	Truck	0
3	4	365	13330	126	147	0.000	SUV	1
4	5	365	22533	10	11	0.001	Truck	0

In [138]: `sim_sum_tot.groupby(['VehicleType'])['Loss'].count()`

Out[138]:

```
VehicleType
Car          9085
Minivan      1520
SUV          7463
Truck       11932
Name: Loss, dtype: int64
```

This dataset contains 30,000 vehicles, of which 4031 have been in a collision. For a further breakdown of how many vehicles have been in a collision by VehicleType:

In [139]: `sim_sum_tot.groupby(['VehicleType', 'Loss'])['Loss'].count()`

Out[139]:

```
VehicleType  Loss
Car          0      7955
             1      1130
Minivan      0      1365
             1       155
SUV          0      6368
             1      1095
Truck        0     10281
             1      1651
Name: Loss, dtype: int64
```

## (4) Is there a statistically significant difference between vehicle types?

If we assume the loss populations are sampled from a binomial distribution with probability given by  $\text{TotalLossPerType}/\text{TotalType}$  a simple z-test can be conducted to determine if the null hypothesis, the distributions being 'sampled' from a similar distribution, can be rejected. For a significance  $\alpha = 0.05$  a z-value greater than the critical value of 1.64 means the null hypothesis is rejected and the population proportions are statistically significantly different at the 0.05 significance level. However, there is a nonzero chance that continually looking at distributions will result in a positive effect (The Look-Elsewhere effect) one way to combat this is to divide the significance value you're looking for by the number of unique trials (here 6) and using that critical value. This then is  $\alpha = 0.083$  and chances  $z_{critical} = 2.64$ . However, the choice of  $p < 0.05$  rejecting the null hypothesis is a convention and the distinction here is somewhat arbitrary. The conclusions that I draw depend on how liberal we want to be in the definition of statistical significance.

$$z = \frac{p_1 - p_2}{\sqrt{p(1-p)(1/n_1 + 1/n_2)}}$$

```
In [140]: ProbDict={}
          ProbDict['Car']      = [1130.,9085.];
          ProbDict['Minivan'] = [155.,1520.] ;
          ProbDict['SUV']     = [1095.,7463.];
          ProbDict['Truck']   = [1651.,11932.];
```

```
In [141]: def calculateZ(ProbDict,Vehicle1,Vehicle2):
    #Takes the total sample size of two distributions and number of 'favorabl
    e' cases here the Loss and returns a z-test value
    n1=ProbDict[Vehicle1][1]
    x1=ProbDict[Vehicle1][0]
    n2=ProbDict[Vehicle2][1]
    x2=ProbDict[Vehicle2][0]
    p1 = x1/n1
    p2 = x2/n2
    p = (x1+x2)/(n1+n2)
    z = np.abs(p1-p2)/np.sqrt(p*(1-p)*(1/n1+1/n2))
    print('z value for %s and %s: %.2f'%(Vehicle1,Vehicle2,z))
    return z
usedKeys=[]
for key in ProbDict:
    usedKeys.append(key)
    for key2 in ProbDict:
        if key !=key2 and key2 not in usedKeys:
            calculateZ(ProbDict,key,key2)
```

```
z value for Car and Minivan: 2.48
z value for Car and SUV: 4.19
z value for Car and Truck: 2.96
z value for Minivan and SUV: 4.59
z value for Minivan and Truck: 3.92
z value for SUV and Truck: 1.62
```

**Applying this z test for all possible combinations all combinations can be said to be statistically significant except for the combination of populations of Trucks and SUVs. Including the Look-elsewhere effect also can bring back into statistical insignificance the combination of Cars/Minivans**

## (5) Are hard brakes and hard accelerations equally important in predicting risks?

```
In [142]: print("Hard Brakes per Loss Event- Mean: %.2f, Stdev: %.2f, Median: %.2f"%(si
m_sum_tot[sim_sum_tot['Loss']==1]['HardBrakes'].mean(),sim_sum_tot[sim_sum_tot
['Loss']==1]['HardBrakes'].std(), sim_sum_tot[sim_sum_tot['Loss']==1]['HardBra
kes'].median()))
print("Hard Accelerations per Loss Event- Mean: %.2f, Stdev: %.2f, Median: %.2
f"%(sim_sum_tot[sim_sum_tot['Loss']==1]['HardAccelerations'].mean(),sim_sum_t
ot[sim_sum_tot['Loss']==1]['HardAccelerations'].std(),sim_sum_tot[sim_sum_tot[
'Loss']==1]['HardAccelerations'].median()))
```

```
Hard Brakes per Loss Event- Mean: 170.24, Stdev: 495.04, Median: 98.00
Hard Accelerations per Loss Event- Mean: 138.25, Stdev: 534.86, Median: 68.00
```


```
In [143]: print("Hard Brakes per 0 Loss Event- Mean: %.2f, Stdev: %.2f, Median: %.2f"%(
sim_sum_tot[sim_sum_tot['Loss']==0]['HardBrakes'].mean(),sim_sum_tot[sim_sum_t
ot['Loss']==0]['HardBrakes'].std(), sim_sum_tot[sim_sum_tot['Loss']==0]['HardB
rakes'].median()))
print("Hard Accelerations per 0 Loss Event- Mean: %.2f, Stdev: %.2f, Median:
%.2f"%(sim_sum_tot[sim_sum_tot['Loss']==0]['HardAccelerations'].mean(),sim_su
m_tot[sim_sum_tot['Loss']==0]['HardAccelerations'].std(),sim_sum_tot[sim_sum_t
ot['Loss']==0]['HardAccelerations'].median()))
```

Hard Brakes per 0 Loss Event- Mean: 167.44, Stdev: 570.00, Median: 98.00

Hard Accelerations per 0 Loss Event- Mean: 104.53, Stdev: 376.43, Median: 56.00

```
In [144]: sim_sum_tot[sim_sum_tot['Loss']==0]['HardAccelerations'].plot()
```

Out[144]: <matplotlib.axes.\_subplots.AxesSubplot at 0x19b81308948>



```
In [145]: sim_sum_tot[sim_sum_tot['Loss']==1].describe()
```

Out[145]:

	Vehicle	Days	Distance	HardBrakes	HardAccelerations	NightTime_Pct	Lo
<b>count</b>	4031.000000	4031.0	4031.000000	4031.000000	4031.000000	4031.000000	4031.000000
<b>mean</b>	15052.420491	365.0	14757.022823	170.241131	138.251054	0.031782	0.031782
<b>std</b>	8614.665175	0.0	4514.713625	495.044167	534.856132	0.054822	0.054822
<b>min</b>	4.000000	365.0	2488.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	7653.500000	365.0	11550.500000	41.000000	25.000000	0.001000	0.001000
<b>50%</b>	15071.000000	365.0	14369.000000	98.000000	68.000000	0.008000	0.008000
<b>75%</b>	22361.000000	365.0	17477.500000	208.000000	159.000000	0.038000	0.038000
<b>max</b>	29981.000000	365.0	34805.000000	16315.000000	24629.000000	0.761000	0.761000

```
In [146]: sim_sum_tot[sim_sum_tot['Loss']==0].describe()
```

Out[146]:

	Vehicle	Days	Distance	HardBrakes	HardAccelerations	NightTime_Pct	
<b>count</b>	25969.000000	25969.0	25969.000000	25969.000000	25969.000000	25969.000000	25969.000000
<b>mean</b>	14992.440718	365.0	13440.090762	167.442951	104.528130	0.029429	
<b>std</b>	8667.612853	0.0	4204.589699	569.999916	376.432155	0.054906	
<b>min</b>	1.000000	365.0	2911.000000	0.000000	0.000000	0.000000	
<b>25%</b>	7477.000000	365.0	10405.000000	39.000000	20.000000	0.001000	
<b>50%</b>	14992.000000	365.0	13039.000000	98.000000	56.000000	0.007000	
<b>75%</b>	22530.000000	365.0	16028.000000	205.000000	129.000000	0.033000	
<b>max</b>	30000.000000	365.0	35159.000000	35639.000000	38221.000000	0.888000	

## Model

This model will be employing a densely connected feed forward neural network for event classification Activation functions: rectified linear unit (ReLU) activation function such that the gradients of activation functions will not diminish as you increase the depth of the network sigmoid activation function is applied to the output layer as is standard for classification problems

Loss/Response Function: Binary Cross Entropy will be used to calculate the predicted probability (probability of correct classification based on the input values) Larger predicted probabilities correspond to lower response function values for correctly identified events.

Optimization Function: Adam (Adaptive moment estimation) Optimizer which is the standard efficient optimization function that calculates adaptive learning rates for all parameters in the network. It computes exponentially weighted averages of past gradients and squares of gradients

I'll split up the input dataset into 2 dataframes, X:the useful input values (dropping Loss as it is the classifier, Vehicle which is effectively an index, and Days which is 365 for every row, and y: the truth value of loss prediction

First I'll show an extremely naive approach just taking in the disparate raw datasets before moving onto using the same approach with oversampled minority case events using both the SMOTE and ADASYN oversampling techniques.

```
In [147]: sim_sum_tot = sim_sum_tot.replace("Car", 0)
sim_sum_tot = sim_sum_tot.replace("SUV", 1)
sim_sum_tot = sim_sum_tot.replace("Truck", 2)
sim_sum_tot = sim_sum_tot.replace("Minivan", 3)

X = sim_sum_tot.drop(['Loss', 'Vehicle', 'Days'], axis=1)
y = sim_sum_tot['Loss']
```

```
In [148]: #splitting up the signal and background datasets into training sets, testing sets, and validation sets.
#This splits into 64% to be used for training and 16% testing and 16% validation
ix = range(y.shape[0])
X_train, X_test, y_train, y_test, ix_train, ix_test = train_test_split(X, y, ix, train_size=0.8)
X_train, X_val, y_train, y_val, ix_train, ix_val = train_test_split(X_train, y_train, ix_train, test_size=0.2)
```

```
In [149]: #Neural networks prefer inputs that are similar to each other while at the same time as normally distributed as possible
#sklearn has a variety of scalers that can be used to perform this transformation, here I use RobustScaler
from sklearn.preprocessing import RobustScaler
scaler = RobustScaler()
X_train = scaler.fit_transform(X_train)
X_val = scaler.transform(X_val)
X_test = scaler.transform(X_test)
```

```
In [150]: #A Function for making various depths and complexities of dense neural networks
def DNNmodel(Input_shape=(10,), n_hidden=1, n_nodesHidden=20, dropout=0.2, optimizer='adam'):
    inputs=Input(shape=Input_shape)
    i=0
    if n_hidden>0:
        hidden=Dense(n_nodesHidden, activation='relu')(inputs)
        hidden=Dropout(dropout)(hidden)
        i+=1
    while i<n_hidden:
        hidden=Dense(n_nodesHidden, activation='relu')(hidden)
        hidden=Dropout(dropout)(hidden)
        i+=1
    outputs = Dense(1,activation='sigmoid')(hidden)
    model = Model(inputs,outputs)
    model.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accuracy'])
    model.summary()
    return model
```

```
In [151]: model = DNNmodel(Input_shape=(5,),n_hidden=2)
```

Model: "model\_15"

Layer (type)	Output Shape	Param #
=====		
input_15 (InputLayer)	(None, 5)	0
dense_56 (Dense)	(None, 20)	120
dropout_42 (Dropout)	(None, 20)	0
dense_57 (Dense)	(None, 20)	420
dropout_43 (Dropout)	(None, 20)	0
dense_58 (Dense)	(None, 1)	21
=====		
Total params: 561		
Trainable params: 561		
Non-trainable params: 0		



```
In [152]: model.fit(X_train,y_train,epochs=50,callbacks=[EarlyStopping(verbose=True,patience=10,monitor='val_loss')], validation_data=(X_val, y_val),batch_size=10)
```

Train on 19200 samples, validate on 4800 samples

Epoch 1/50

19200/19200 [=====] - 2s 122us/step - loss: 0.4184 - accuracy: 0.8645 - val\_loss: 0.4054 - val\_accuracy: 0.8610

Epoch 2/50

19200/19200 [=====] - 2s 105us/step - loss: 0.3952 - accuracy: 0.8670 - val\_loss: 0.4040 - val\_accuracy: 0.8606

Epoch 3/50

19200/19200 [=====] - 2s 107us/step - loss: 0.3906 - accuracy: 0.8672 - val\_loss: 0.4031 - val\_accuracy: 0.8610

Epoch 4/50

19200/19200 [=====] - 2s 106us/step - loss: 0.3893 - accuracy: 0.8672 - val\_loss: 0.4022 - val\_accuracy: 0.8610

Epoch 5/50

19200/19200 [=====] - 2s 106us/step - loss: 0.3897 - accuracy: 0.8672 - val\_loss: 0.4022 - val\_accuracy: 0.8610

Epoch 6/50

19200/19200 [=====] - 2s 105us/step - loss: 0.3884 - accuracy: 0.8672 - val\_loss: 0.4014 - val\_accuracy: 0.8610

Epoch 7/50

19200/19200 [=====] - 2s 104us/step - loss: 0.3892 - accuracy: 0.8671 - val\_loss: 0.4002 - val\_accuracy: 0.8610

Epoch 8/50

19200/19200 [=====] - 2s 105us/step - loss: 0.3883 - accuracy: 0.8672 - val\_loss: 0.3993 - val\_accuracy: 0.8612

Epoch 9/50

19200/19200 [=====] - 2s 105us/step - loss: 0.3883 - accuracy: 0.8672 - val\_loss: 0.3999 - val\_accuracy: 0.8610

Epoch 10/50

19200/19200 [=====] - 2s 105us/step - loss: 0.3876 - accuracy: 0.8671 - val\_loss: 0.3999 - val\_accuracy: 0.8610

Epoch 11/50

19200/19200 [=====] - 2s 106us/step - loss: 0.3882 - accuracy: 0.8672 - val\_loss: 0.3989 - val\_accuracy: 0.8610

Epoch 12/50

19200/19200 [=====] - 2s 105us/step - loss: 0.3866 - accuracy: 0.8672 - val\_loss: 0.3992 - val\_accuracy: 0.8610

Epoch 13/50

19200/19200 [=====] - 2s 104us/step - loss: 0.3869 - accuracy: 0.8672 - val\_loss: 0.3988 - val\_accuracy: 0.8610

Epoch 14/50

19200/19200 [=====] - 2s 110us/step - loss: 0.3864 - accuracy: 0.8671 - val\_loss: 0.4002 - val\_accuracy: 0.8610

Epoch 15/50

19200/19200 [=====] - 2s 111us/step - loss: 0.3857 - accuracy: 0.8672 - val\_loss: 0.3997 - val\_accuracy: 0.8608

Epoch 16/50

19200/19200 [=====] - 2s 108us/step - loss: 0.3856 - accuracy: 0.8672 - val\_loss: 0.3993 - val\_accuracy: 0.8610

Epoch 17/50

19200/19200 [=====] - 2s 108us/step - loss: 0.3857 - accuracy: 0.8672 - val\_loss: 0.3997 - val\_accuracy: 0.8610

Epoch 18/50

19200/19200 [=====] - 2s 109us/step - loss: 0.3853 - accuracy: 0.8672 - val\_loss: 0.3993 - val\_accuracy: 0.8610

Epoch 19/50

19200/19200 [=====] - 2s 107us/step - loss: 0.3862 -

```
accuracy: 0.8672 - val_loss: 0.3999 - val_accuracy: 0.8610
Epoch 20/50
19200/19200 [=====] - 2s 107us/step - loss: 0.3856 -
accuracy: 0.8671 - val_loss: 0.3993 - val_accuracy: 0.8610
Epoch 21/50
19200/19200 [=====] - 2s 107us/step - loss: 0.3851 -
accuracy: 0.8671 - val_loss: 0.3987 - val_accuracy: 0.8610
Epoch 22/50
19200/19200 [=====] - 2s 105us/step - loss: 0.3856 -
accuracy: 0.8672 - val_loss: 0.3999 - val_accuracy: 0.8610
Epoch 23/50
19200/19200 [=====] - 2s 106us/step - loss: 0.3856 -
accuracy: 0.8672 - val_loss: 0.3992 - val_accuracy: 0.8610
Epoch 24/50
19200/19200 [=====] - 2s 105us/step - loss: 0.3854 -
accuracy: 0.8673 - val_loss: 0.3988 - val_accuracy: 0.8610
Epoch 25/50
19200/19200 [=====] - 2s 106us/step - loss: 0.3854 -
accuracy: 0.8672 - val_loss: 0.3991 - val_accuracy: 0.8610
Epoch 26/50
19200/19200 [=====] - 2s 105us/step - loss: 0.3849 -
accuracy: 0.8672 - val_loss: 0.3992 - val_accuracy: 0.8610
Epoch 27/50
19200/19200 [=====] - 2s 106us/step - loss: 0.3848 -
accuracy: 0.8672 - val_loss: 0.3991 - val_accuracy: 0.8610
Epoch 28/50
19200/19200 [=====] - 2s 105us/step - loss: 0.3855 -
accuracy: 0.8672 - val_loss: 0.3991 - val_accuracy: 0.8610
Epoch 29/50
19200/19200 [=====] - 2s 106us/step - loss: 0.3853 -
accuracy: 0.8672 - val_loss: 0.3999 - val_accuracy: 0.8610
Epoch 30/50
19200/19200 [=====] - 2s 106us/step - loss: 0.3851 -
accuracy: 0.8673 - val_loss: 0.3986 - val_accuracy: 0.8610
Epoch 31/50
19200/19200 [=====] - 2s 107us/step - loss: 0.3845 -
accuracy: 0.8671 - val_loss: 0.3994 - val_accuracy: 0.8610
Epoch 32/50
19200/19200 [=====] - 2s 104us/step - loss: 0.3851 -
accuracy: 0.8672 - val_loss: 0.4008 - val_accuracy: 0.8610
Epoch 33/50
19200/19200 [=====] - 2s 105us/step - loss: 0.3848 -
accuracy: 0.8672 - val_loss: 0.3991 - val_accuracy: 0.8610
Epoch 34/50
19200/19200 [=====] - 2s 103us/step - loss: 0.3843 -
accuracy: 0.8673 - val_loss: 0.3995 - val_accuracy: 0.8610
Epoch 35/50
19200/19200 [=====] - 2s 104us/step - loss: 0.3852 -
accuracy: 0.8672 - val_loss: 0.3992 - val_accuracy: 0.8610
Epoch 36/50
19200/19200 [=====] - 2s 109us/step - loss: 0.3842 -
accuracy: 0.8673 - val_loss: 0.3989 - val_accuracy: 0.8610
Epoch 37/50
19200/19200 [=====] - 2s 107us/step - loss: 0.3846 -
accuracy: 0.8671 - val_loss: 0.3994 - val_accuracy: 0.8610
Epoch 38/50
19200/19200 [=====] - 2s 105us/step - loss: 0.3840 -
```

```

accuracy: 0.8672 - val_loss: 0.4000 - val_accuracy: 0.8610
Epoch 39/50
19200/19200 [=====] - 2s 106us/step - loss: 0.3846 -
accuracy: 0.8672 - val_loss: 0.3993 - val_accuracy: 0.8610
Epoch 40/50
19200/19200 [=====] - 2s 106us/step - loss: 0.3845 -
accuracy: 0.8673 - val_loss: 0.3991 - val_accuracy: 0.8610
Epoch 00040: early stopping

```

Out[152]: <keras.callbacks.callbacks.History at 0x19b82184fc8>

```

In [153]: history = model.history.history
print("history keys: ", history.keys())

```


history keys: dict\_keys(['val\_loss', 'val\_accuracy', 'loss', 'accuracy'])

```

In [154]: #Accuracy plot
plt.plot(100 * np.array(history['accuracy']), label='training')
plt.plot(100 * np.array(history['val_accuracy']), label='validation')
plt.xlim(0)
plt.xlabel('epoch')
plt.ylabel('accuracy %')
plt.legend(loc='lower right', fontsize=20)

```

Out[154]: <matplotlib.legend.Legend at 0x19b82504dc8>



```

In [155]: yhat = model.predict(X_test, verbose = True, batch_size = 512)

```

6000/6000 [=====] - 0s 52us/step


```

In [156]: yhat_cls = np.argmax(yhat, axis=1)

```

```
In [157]: bins = np.linspace(0, 1, 20)
_ = plt.hist(yhat[y_test==1], histtype='stepfilled', alpha=0.5, color='red', label=r"Loss", bins=bins)
_ = plt.hist(yhat[y_test==0], histtype='stepfilled', alpha=0.5, color='blue', label=r"NoLoss", bins=bins)
plt.axvline(x=0.5, linestyle='--')
plt.legend(loc='upper right', fontsize=20)
plt.xlabel('P(Loss Event)', fontsize='large')
plt.ylabel('Events', fontsize='large')
```

Out[157]: Text(0, 0.5, 'Events')



```
In [158]: LossAcc = yhat[y_test==1].round().sum()/y_test.sum()
print("Loss Event Accuracy: %.1f%%" %(LossAcc*100))
```

Loss Event Accuracy: 0.0%

It is obvious from this that the neural network, at least with the sample given is not sufficient without taking into consideration the disparate nature of the signal/background like events. Even though there is clearly a shape difference in the distributions all events are classified as 'No Loss' because of the large pull from the significant number of 'No Loss' events compared to 'Loss' Events. The network accuracy goes to the percentage of events in the majority case while classifying all signal events incorrectly. To account for this I will preform oversampling

## Resampling, NN part 2

Let's try to do some resampling using Adaptive Synthetic Sampling, focuses on data that are hard to learn, fills in more edge cases than the next technique (SMOTE) which fills more linearly between events. [https://imbalanced-learn.readthedocs.io/en/stable/generated/imblearn.over\\_sampling.ADASYN.html](https://imbalanced-learn.readthedocs.io/en/stable/generated/imblearn.over_sampling.ADASYN.html) ([https://imbalanced-learn.readthedocs.io/en/stable/generated/imblearn.over\\_sampling.ADASYN.html](https://imbalanced-learn.readthedocs.io/en/stable/generated/imblearn.over_sampling.ADASYN.html))

```
In [87]: #Upscale the minority class such that each class has roughly equal representat  
ion  
from imblearn.over_sampling import ADASYN  
adasyn = ADASYN(sampling_strategy="auto")  
X_train_adasyn, y_train_adasyn = adasyn.fit_resample(X_train, y_train)  
np.bincount(y_train_adasyn)
```

```
Out[87]: array([16643, 17066], dtype=int64)
```

```
In [88]: model = DNNmodel(Input_shape=(5,),n_hidden=3);  
model.fit(X_train_adasyn,y_train_adasyn,epochs=100,callbacks=[EarlyStopping(verbose=True,patience=20,monitor='val_loss')], validation_data=(X_val, y_val),batch_size=10)
```

Model: "model\_12"

Layer (type)	Output Shape	Param #
input_12 (InputLayer)	(None, 5)	0
dense_44 (Dense)	(None, 20)	120
dropout_33 (Dropout)	(None, 20)	0
dense_45 (Dense)	(None, 20)	420
dropout_34 (Dropout)	(None, 20)	0
dense_46 (Dense)	(None, 20)	420
dropout_35 (Dropout)	(None, 20)	0
dense_47 (Dense)	(None, 1)	21

=====  
Total params: 981

Trainable params: 981

Non-trainable params: 0

=====  
Train on 33709 samples, validate on 4800 samples

Epoch 1/100

33709/33709 [=====] - 4s 111us/step - loss: 0.6916 - accuracy: 0.5394 - val\_loss: 0.7005 - val\_accuracy: 0.4908

Epoch 2/100

33709/33709 [=====] - 3s 103us/step - loss: 0.6868 - accuracy: 0.5501 - val\_loss: 0.7090 - val\_accuracy: 0.5048

Epoch 3/100

33709/33709 [=====] - 4s 105us/step - loss: 0.6848 - accuracy: 0.5558 - val\_loss: 0.6963 - val\_accuracy: 0.5117

Epoch 4/100

33709/33709 [=====] - 3s 101us/step - loss: 0.6841 - accuracy: 0.5615 - val\_loss: 0.7075 - val\_accuracy: 0.4998

Epoch 5/100

33709/33709 [=====] - 3s 102us/step - loss: 0.6829 - accuracy: 0.5615 - val\_loss: 0.6887 - val\_accuracy: 0.5098

Epoch 6/100

33709/33709 [=====] - 3s 103us/step - loss: 0.6820 - accuracy: 0.5625 - val\_loss: 0.6879 - val\_accuracy: 0.5081

Epoch 7/100

33709/33709 [=====] - 4s 118us/step - loss: 0.6814 - accuracy: 0.5640 - val\_loss: 0.7086 - val\_accuracy: 0.4840

Epoch 8/100

33709/33709 [=====] - 4s 132us/step - loss: 0.6807 - accuracy: 0.5656 - val\_loss: 0.6937 - val\_accuracy: 0.5106

Epoch 9/100

33709/33709 [=====] - 4s 116us/step - loss: 0.6803 - accuracy: 0.5686 - val\_loss: 0.7089 - val\_accuracy: 0.5227

Epoch 10/100

33709/33709 [=====] - 3s 101us/step - loss: 0.6798 - accuracy: 0.5663 - val\_loss: 0.6627 - val\_accuracy: 0.6158

Epoch 11/100

33709/33709 [=====] - 4s 116us/step - loss: 0.6800 -



```
accuracy: 0.5673 - val_loss: 0.6946 - val_accuracy: 0.5475
Epoch 12/100
33709/33709 [=====] - 3s 99us/step - loss: 0.6782 -
accuracy: 0.5695 - val_loss: 0.6664 - val_accuracy: 0.6025
Epoch 13/100
33709/33709 [=====] - 3s 98us/step - loss: 0.6794 -
accuracy: 0.5709 - val_loss: 0.6668 - val_accuracy: 0.6219
Epoch 14/100
33709/33709 [=====] - 3s 101us/step - loss: 0.6781 -
accuracy: 0.5694 - val_loss: 0.7003 - val_accuracy: 0.5362
Epoch 15/100
33709/33709 [=====] - 3s 100us/step - loss: 0.6787 -
accuracy: 0.5709 - val_loss: 0.7060 - val_accuracy: 0.5581
Epoch 16/100
33709/33709 [=====] - 3s 103us/step - loss: 0.6774 -
accuracy: 0.5739 - val_loss: 0.6679 - val_accuracy: 0.5942
Epoch 17/100
33709/33709 [=====] - 3s 102us/step - loss: 0.6775 -
accuracy: 0.5758 - val_loss: 0.6853 - val_accuracy: 0.5490
Epoch 18/100
33709/33709 [=====] - 3s 103us/step - loss: 0.6762 -
accuracy: 0.5751 - val_loss: 0.6858 - val_accuracy: 0.5540
Epoch 19/100
33709/33709 [=====] - 3s 103us/step - loss: 0.6766 -
accuracy: 0.5759 - val_loss: 0.7008 - val_accuracy: 0.5437
Epoch 20/100
33709/33709 [=====] - 3s 101us/step - loss: 0.6753 -
accuracy: 0.5737 - val_loss: 0.6863 - val_accuracy: 0.5500
Epoch 21/100
33709/33709 [=====] - 3s 101us/step - loss: 0.6760 -
accuracy: 0.5742 - val_loss: 0.6751 - val_accuracy: 0.5804
Epoch 22/100
33709/33709 [=====] - 3s 100us/step - loss: 0.6756 -
accuracy: 0.5755 - val_loss: 0.7007 - val_accuracy: 0.5408
Epoch 23/100
33709/33709 [=====] - 4s 106us/step - loss: 0.6748 -
accuracy: 0.5784 - val_loss: 0.7159 - val_accuracy: 0.5156
Epoch 24/100
33709/33709 [=====] - 3s 103us/step - loss: 0.6763 -
accuracy: 0.5758 - val_loss: 0.6906 - val_accuracy: 0.5598
Epoch 25/100
33709/33709 [=====] - 3s 102us/step - loss: 0.6768 -
accuracy: 0.5726 - val_loss: 0.6785 - val_accuracy: 0.5694
Epoch 26/100
33709/33709 [=====] - 3s 101us/step - loss: 0.6758 -
accuracy: 0.5758 - val_loss: 0.6845 - val_accuracy: 0.5481
Epoch 27/100
33709/33709 [=====] - 3s 100us/step - loss: 0.6764 -
accuracy: 0.5759 - val_loss: 0.6971 - val_accuracy: 0.5510
Epoch 28/100
33709/33709 [=====] - 4s 107us/step - loss: 0.6756 -
accuracy: 0.5768 - val_loss: 0.6980 - val_accuracy: 0.5537
Epoch 29/100
33709/33709 [=====] - 3s 104us/step - loss: 0.6750 -
accuracy: 0.5772 - val_loss: 0.7004 - val_accuracy: 0.5269
Epoch 30/100
33709/33709 [=====] - 3s 101us/step - loss: 0.6754 -
```


accuracy: 0.5779 - val\_loss: 0.7093 - val\_accuracy: 0.5421  
Epoch 00030: early stopping

Out[88]: <keras.callbacks.callbacks.History at 0x19bfb37fe08>

```
In [89]: history = model.history.history
print("history keys: ", history.keys())
#Accuracy plot
plt.plot(100 * np.array(history['accuracy']), label='training')
plt.plot(100 * np.array(history['val_accuracy']), label='validation')
plt.xlim(0)
plt.xlabel('epoch')
plt.ylabel('accuracy %')
plt.legend(loc='lower right', fontsize=20)
```

history keys: dict\_keys(['val\_loss', 'val\_accuracy', 'loss', 'accuracy'])

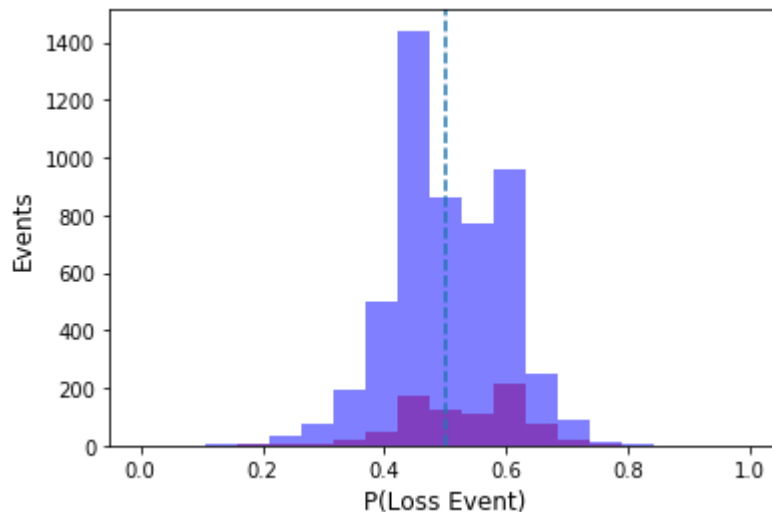
Out[89]: <matplotlib.legend.Legend at 0x19bfc9dbf48>



```
In [90]: yhat = model.predict(X_test, verbose = True, batch_size = 512)
yhat_cls = np.argmax(yhat, axis=1)
sum(yhat_cls)
bins = np.linspace(0, 1, 20)
_ = plt.hist(yhat[y_test==1], histtype='stepfilled', alpha=0.5, color='red', label=r"Signal", bins=bins)
_ = plt.hist(yhat[y_test==0], histtype='stepfilled', alpha=0.5, color='blue', label=r'Background', bins=bins)
plt.axvline(x=0.5, linestyle='--')
plt.xlabel('P(Loss Event)', fontsize='large')
plt.ylabel('Events', fontsize='large')
```

6000/6000 [=====] - 0s 44us/step

Out[90]: Text(0, 0.5, 'Events')



```
In [91]: print("Loss Events: mean: %.3f, std: %.3f" %(yhat[y_test==1].mean(),yhat[y_test==1].std()))
print("NoLoss Events: mean: %.3f, std: %.3f" %(yhat[y_test==0].mean(),yhat[y_test==0].std()))
```

Loss Events: mean: 0.535, std: 0.094  
NoLoss Events: mean: 0.503, std: 0.092

```
In [92]: LossAcc = yhat[y_test==1].round().sum()/y_test.sum()
print("Loss Event Accuracy: %.1f%" %(LossAcc*100))
```

Loss Event Accuracy: 61.2%

## Resampling using SMOTE

Synthetic minority oversampling technique (Smote) generates synthetic data that is similar to, but not exactly like the minority class using nearest-neighbors approach to fill in space between neighbors.

```
In [99]: from imblearn.over_sampling import SMOTE
smote = SMOTE(sampling_strategy="auto")
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)
np.bincount(y_train_smote)
```

```
Out[99]: array([16643, 16643], dtype=int64)
```

```
In [100]: model = DNNmodel(Input_shape=(5,),n_hidden=3);  
model.fit(X_train_smote,y_train_smote,epochs=100,callbacks=[EarlyStopping(verbose=True,patience=20,monitor='val_loss')], validation_data=(X_val, y_val),batch_size=10)
```

Model: "model\_14"

Layer (type)	Output Shape	Param #
input_14 (InputLayer)	(None, 5)	0
dense_52 (Dense)	(None, 20)	120
dropout_39 (Dropout)	(None, 20)	0
dense_53 (Dense)	(None, 20)	420
dropout_40 (Dropout)	(None, 20)	0
dense_54 (Dense)	(None, 20)	420
dropout_41 (Dropout)	(None, 20)	0
dense_55 (Dense)	(None, 1)	21

=====  
Total params: 981

Trainable params: 981

Non-trainable params: 0

=====  
Train on 33286 samples, validate on 4800 samples

Epoch 1/100

33286/33286 [=====] - 4s 116us/step - loss: 0.6892 - accuracy: 0.5508 - val\_loss: 0.6951 - val\_accuracy: 0.5235

Epoch 2/100

33286/33286 [=====] - 4s 106us/step - loss: 0.6828 - accuracy: 0.5616 - val\_loss: 0.6788 - val\_accuracy: 0.5631

Epoch 3/100

33286/33286 [=====] - 4s 106us/step - loss: 0.6830 - accuracy: 0.5649 - val\_loss: 0.6618 - val\_accuracy: 0.6058

Epoch 4/100

33286/33286 [=====] - 3s 103us/step - loss: 0.6809 - accuracy: 0.5674 - val\_loss: 0.6851 - val\_accuracy: 0.5702

Epoch 5/100

33286/33286 [=====] - 3s 103us/step - loss: 0.6804 - accuracy: 0.5690 - val\_loss: 0.6716 - val\_accuracy: 0.5754

Epoch 6/100

33286/33286 [=====] - 3s 105us/step - loss: 0.6791 - accuracy: 0.5714 - val\_loss: 0.6953 - val\_accuracy: 0.5331

Epoch 7/100

33286/33286 [=====] - 3s 103us/step - loss: 0.6779 - accuracy: 0.5724 - val\_loss: 0.6850 - val\_accuracy: 0.5369

Epoch 8/100

33286/33286 [=====] - 3s 103us/step - loss: 0.6792 - accuracy: 0.5745 - val\_loss: 0.6838 - val\_accuracy: 0.5517

Epoch 9/100

33286/33286 [=====] - 3s 104us/step - loss: 0.6784 - accuracy: 0.5749 - val\_loss: 0.6714 - val\_accuracy: 0.5775

Epoch 10/100

33286/33286 [=====] - 3s 103us/step - loss: 0.6770 - accuracy: 0.5784 - val\_loss: 0.6645 - val\_accuracy: 0.5925

Epoch 11/100

33286/33286 [=====] - 3s 103us/step - loss: 0.6760 -


```
accuracy: 0.5787 - val_loss: 0.6794 - val_accuracy: 0.5717
Epoch 12/100
33286/33286 [=====] - 3s 104us/step - loss: 0.6763 -
accuracy: 0.5771 - val_loss: 0.6724 - val_accuracy: 0.5738
Epoch 13/100
33286/33286 [=====] - 3s 102us/step - loss: 0.6759 -
accuracy: 0.5777 - val_loss: 0.6974 - val_accuracy: 0.5419
Epoch 14/100
33286/33286 [=====] - 3s 103us/step - loss: 0.6753 -
accuracy: 0.5759 - val_loss: 0.6897 - val_accuracy: 0.5606
Epoch 15/100
33286/33286 [=====] - 3s 104us/step - loss: 0.6750 -
accuracy: 0.5776 - val_loss: 0.6942 - val_accuracy: 0.5546
Epoch 16/100
33286/33286 [=====] - 3s 102us/step - loss: 0.6747 -
accuracy: 0.5790 - val_loss: 0.6757 - val_accuracy: 0.5606
Epoch 17/100
33286/33286 [=====] - 3s 103us/step - loss: 0.6752 -
accuracy: 0.5802 - val_loss: 0.6746 - val_accuracy: 0.5817
Epoch 18/100
33286/33286 [=====] - 3s 103us/step - loss: 0.6739 -
accuracy: 0.5805 - val_loss: 0.6684 - val_accuracy: 0.5810
Epoch 19/100
33286/33286 [=====] - 3s 102us/step - loss: 0.6729 -
accuracy: 0.5822 - val_loss: 0.6651 - val_accuracy: 0.6002
Epoch 20/100
33286/33286 [=====] - 4s 108us/step - loss: 0.6737 -
accuracy: 0.5810 - val_loss: 0.6740 - val_accuracy: 0.5967
Epoch 21/100
33286/33286 [=====] - 4s 107us/step - loss: 0.6724 -
accuracy: 0.5827 - val_loss: 0.6816 - val_accuracy: 0.5863
Epoch 22/100
33286/33286 [=====] - 3s 101us/step - loss: 0.6730 -
accuracy: 0.5806 - val_loss: 0.6693 - val_accuracy: 0.5994
Epoch 23/100
33286/33286 [=====] - 3s 102us/step - loss: 0.6743 -
accuracy: 0.5829 - val_loss: 0.6756 - val_accuracy: 0.5700
Epoch 00023: early stopping
```

Out[100]: <keras.callbacks.callbacks.History at 0x19bfc588>

```
In [101]: history = model.history.history
print("history keys: ", history.keys())
#Accuracy plot
plt.plot(100 * np.array(history['accuracy']), label='training')
plt.plot(100 * np.array(history['val_accuracy']), label='validation')
plt.xlim(0)
plt.xlabel('epoch')
plt.ylabel('accuracy %')
plt.legend(loc='lower right', fontsize=20)
```

history keys: dict\_keys(['val\_loss', 'val\_accuracy', 'loss', 'accuracy'])

Out[101]: <matplotlib.legend.Legend at 0x19bfe356a08>






```
In [102]: yhat = model.predict(X_test, verbose = True, batch_size = 512)
bins = np.linspace(0, 1, 20)
_ = plt.hist(yhat[y_test==1], histtype='stepfilled', alpha=0.5, color='red', label=r"Signal", bins=bins)
_ = plt.hist(yhat[y_test==0], histtype='stepfilled', alpha=0.5, color='blue', label=r'Background', bins=bins)
plt.axvline(x=0.5, linestyle='--')
plt.xlabel('P(Loss Event)', fontsize='large')
plt.ylabel('Events', fontsize='large')
```

6000/6000 [=====] - 0s 50us/step

Out[102]: Text(0, 0.5, 'Events')



```
In [103]: print("Loss Events: mean: %.3f, std: %.3f" %(yhat[y_test==1].mean(),yhat[y_test==1].std()))
print("NoLoss Events: mean: %.3f, std: %.3f" %(yhat[y_test==0].mean(),yhat[y_test==0].std()))
```

Loss Events: mean: 0.510, std: 0.096  
NoLoss Events: mean: 0.480, std: 0.097

```
In [104]: LossAcc = yhat[y_test==1].round().sum()/y_test.sum()
print("Loss Event Accuracy: %.1f%%" %(LossAcc*100))
```

Loss Event Accuracy: 55.7%

## Let's try a BDT instead of a NN, with ADASYN Resampling

Change to XGBoost friendlier format, using all potentially useful columns, dropping Vehicle (an index) and Days (always 365, no chance of separation between loss type events)

```
In [105]: #splitting up the signal and background datasets into training sets, testing sets, and validation sets.
#This splits into 80% to be used for training and 20% testing
X = sim_sum_tot.drop(['Vehicle', 'Days'], axis=1)
y = sim_sum_tot['Loss']
ix = range(y.shape[0])
#X=X.drop('Loss',axis=1)
X_train, X_test, y_train, y_test, ix_train, ix_test = train_test_split(X, y, ix, train_size=0.8)
#X_train, X_val, y_train, y_val, ix_train, ix_val=train_test_split(X_train, y_train, ix_train, test_size=0.2)
```

```
In [106]: X.columns
```

```
Out[106]: Index(['Distance', 'HardBrakes', 'HardAccelerations', 'NightTime_Pct', 'VehicleType', 'Loss'],
              dtype='object')
```

```
In [107]: adasyn = ADASYN(sampling_strategy="auto")
X_train_adasyn, y_train_adasyn = adasyn.fit_resample(X_train, y_train)
np.bincount(y_train_adasyn)
```

```
Out[107]: array([20744, 20434], dtype=int64)
```

```
In [108]: X_train_adasyn
```

```
Out[108]:
```

	Distance	HardBrakes	HardAccelerations	NightTime_Pct	VehicleType	Loss
0	14762	541	12	0.191000	3	0
1	13565	113	54	0.000000	0	0
2	25507	132	119	0.005000	0	0
3	21419	454	78	0.000000	1	0
4	8054	57	129	0.057000	0	0
...	...	...	...	...	...	...
41173	11563	72	73	0.042493	1	1
41174	11577	97	64	0.090249	1	1
41175	11558	48	106	0.078863	0	1
41176	11559	49	106	0.075842	0	1
41177	11575	93	69	0.084200	1	1

41178 rows × 6 columns

```
In [109]: print('Number of training samples: {}'.format(len(X_train)))
print('Number of testing samples: {}'.format(len(X_test)))
print('ADASYN: Number of training samples: {}'.format(len(X_train_adasyn)))

print('\nNumber of signal events in training set: {}'.format(len(X_train[X_train.Loss == 1])))
print('ADASYN: Number of signal events in training set: {}'.format(len(X_train_adasyn[X_train_adasyn.Loss == 1])))
print('Number of background events in training set: {}'.format(len(X_train[X_train.Loss == 0])))
print('Fraction signal: {}'.format(len(X_train[X_train.Loss == 1])/(float)(len(X_train[X_train.Loss == 1]) + len(X_train[X_train.Loss == 0]))))
print('ADASYN: Fraction signal: {}'.format(len(X_train_adasyn[X_train_adasyn.Loss == 1])/(float)(len(X_train_adasyn[X_train_adasyn.Loss == 1]) + len(X_train_adasyn[X_train_adasyn.Loss == 0]))))
```

Number of training samples: 24000

Number of testing samples: 6000

ADASYN: Number of training samples: 41178

Number of signal events in training set: 3256

ADASYN: Number of signal events in training set: 20434

Number of background events in training set: 20744

Fraction signal: 0.13566666666666666

ADASYN: Fraction signal: 0.49623585409684784

```
In [110]: features = X_train.columns[:-1] # we skip the last column because it is the Loss label
features
```

```
Out[110]: Index(['Distance', 'HardBrakes', 'HardAccelerations', 'NightTime_Pct',
                'VehicleType'],
                dtype='object')
```

```
In [111]: binary_bdt_param = {
    "learning_rate" : 0.15,
    "max_depth" : 6,
    "colsample_bytree" : 1.0,
    "subsample" : 1.0,
    "n_estimators" : 200,
    "feature_names" : features,
    'objective' : 'binary:logistic' # objective function
}
binary_task_param = {
    "eval_metric" : ["logloss", "error"],
    "early_stopping_rounds" : 30,
    "eval_set": [(X_train_adasyn[features], y_train_adasyn),
                  (X_test[features], y_test)]
}

binary_bdt = xgb.XGBClassifier(**binary_bdt_param)
binary_bdt.fit(X_train_adasyn[features], y_train_adasyn,
               verbose=False, **binary_task_param)
```

[20:07:03] WARNING: C:\Users\Administrator\workspace\xgboost-win64\_release\_1.1.0\src\learner.cc:480:  
Parameters: { feature\_names } might not be used.

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

```
Out[111]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
    colsample_bynode=1, colsample_bytree=1.0,
    feature_names=Index(['Distance', 'HardBrakes', 'HardAcceleration', 'NightTime_Pct',
    'VehicleType'],
    dtype='object'),
    gamma=0, gpu_id=-1, importance_type='gain',
    interaction_constraints='', learning_rate=0.15, max_delta_step=0,
    max_depth=6, min_child_weight=1, missing=nan,
    monotone_constraints='()', n_estimators=200, n_jobs=0,
    num_parallel_tree=1, random_state=0, reg_alpha=0, reg_lambda=1,
    scale_pos_weight=1, subsample=1.0, tree_method='exact',
    validate_parameters=1, verbosity=None)
```

```
In [112]: evaluated_df = X_test.copy()
    evaluated_df["binary_prob"] = binary_bdt.predict_proba(X_test[features])[:,1]
    print(binary_bdt.score(X_test[features], y_test))

    binary_bdt.predict_proba(X_test[features])[:,1].round().sum()
```

0.808

```
Out[112]: 547.0
```

```
In [113]: y_pred = binary_bdt.predict(X_test[features])
          y_pred.sum()
```

```
Out[113]: 547
```

```
In [114]: binary_bdt.predict(X_test[features])
```

```
Out[114]: array([1, 0, 1, ..., 0, 0, 0], dtype=int64)
```

```
In [115]: from sklearn.metrics import accuracy_score
          from sklearn.feature_selection import SelectFromModel

          predictions = binary_bdt.predict(X_test[features])
          accuracy = accuracy_score(y_test, predictions)
          print("Accuracy: %.2f%%" % (accuracy * 100.0))
          thresholds = np.sort(binary_bdt.feature_importances_)
          for thresh in thresholds:
              # select features using threshold
              selection = SelectFromModel(binary_bdt, threshold=thresh, prefit=True)
              select_X_train = selection.transform(X_train_adasyn[features])
              # train model
              selection_model = xgb.XGBClassifier()
              selection_model.fit(select_X_train, y_train_adasyn)
              # eval model
              select_X_test = selection.transform(X_test[features])
              predictions = selection_model.predict(select_X_test)
              accuracy = accuracy_score(y_test, predictions)
              print("Loss Event Accuracy: %.2f%%"%(predictions[y_test==1].sum()/y_test.sum()*100.))
              print("Thresh=%.3f, n=%d, Accuracy: %.2f%%" % (thresh, select_X_train.shape[1], accuracy*100.0))
```

Accuracy: 80.80%

Loss Event Accuracy: 10.32%

Thresh=0.056, n=5, Accuracy: 80.98%

Loss Event Accuracy: 10.32%

Thresh=0.057, n=4, Accuracy: 81.93%

Loss Event Accuracy: 9.16%

Thresh=0.078, n=3, Accuracy: 82.18%


Loss Event Accuracy: 4.77%

Thresh=0.292, n=2, Accuracy: 83.87%

Loss Event Accuracy: 2.97%

Thresh=0.515, n=1, Accuracy: 83.83%



```
In [116]: fig, ax_enum = plt.subplots(1,2, figsize=(16,4))
xgb.plot_importance(binary_bdt, importance_type="weight", ax=ax_enum[0], title=
"Weight", show_values=False, grid=False)
xgb.plot_importance(binary_bdt, importance_type="gain", ax=ax_enum[1], title=
"Gain", show_values=False, grid=False)
plt.ylabel("")
plt.sca(ax_enum[1])
plt.ylabel("")
plt.subplots_adjust(wspace=0.3)
```




importance\_type=Gain is how useful is the variable in terms of separation where importance\_type=Weight is how frequently splitting on the variable occurs

```
In [117]: # plot all predictions (both signal and background)
predictions = binary_bdt.predict_proba(X_test[features])[:,1]
plt.figure();
plt.hist(predictions,bins=np.linspace(0,1,50),histtype='step',color='darkgreen',label='All events');
# make the plot readable
plt.xlabel('Prediction from BDT',fontsize=12);
plt.ylabel('Events',fontsize=12);
plt.legend(frameon=False);

# plot signal and background separately
plt.figure();
plt.hist(binary_bdt.predict_proba(X_test[features][y_test==1])[:,1],bins=np.linspace(0,1,50),
         histtype='step',color='midnightblue',label='Loss');
plt.hist(binary_bdt.predict_proba(X_test[features][y_test==0])[:,1],bins=np.linspace(0,1,50),
         histtype='step',color='firebrick',label='No Loss');
# make the plot readable
plt.xlabel('Prediction from BDT',fontsize=12);
plt.ylabel('Events',fontsize=12);
plt.legend(frameon=False);
```



```
In [118]: plt.figure();  
plt.hist(X_train.Distance[X_train.Loss == 1],bins=np.linspace(0,40000,50),  
         histtype='step',color='midnightblue',label='Loss');  
plt.hist(X_train.Distance[X_train.Loss == 0],bins=np.linspace(0,40000,50),  
         histtype='step',color='firebrick',label='No Loss');  
  
plt.xlabel('Distance',fontsize=12);  
plt.ylabel('Events',fontsize=12);  
plt.legend(frameon=False);
```






```
In [119]: plt.figure();
#plt.plot(X_train_adasyn.Distance[X_train_adasyn.Loss == 0],X_train_adasyn.NightTime_Pct[X_train_adasyn.Loss == 0],
#         'o',markersize=2,color='firebrick',markeredgewidth=0,alpha=0.8,label
#         ='NoLoss');
#plt.plot(X_train_adasyn.Distance[X_train_adasyn.Loss == 1],X_train_adasyn.NightTime_Pct[X_train_adasyn.Loss == 1],
#         'o',markersize=2,color='mediumblue',markeredgewidth=0,alpha=0.8,label
#         ='Loss');

plt.plot(X_train.Distance[X_train.Loss == 0],X_train.NightTime_Pct[X_train.Loss == 0],
         'o',markersize=2,color='firebrick',markeredgewidth=0,alpha=0.8,label=
         'NoLoss');
plt.plot(X_train.Distance[X_train.Loss == 1],X_train.NightTime_Pct[X_train.Loss == 1],
         'o',markersize=2,color='mediumblue',markeredgewidth=0,alpha=0.8,label
         ='Loss');

plt.xlim(0,35000);
plt.ylim(0.0,1);
#plt.yscale('log')
plt.xlabel('Distance',fontsize=12);
plt.ylabel('NightTime_Pct',fontsize=12);
plt.legend(frameon=False,numpoints=1,markerscale=2);
```



## BDT with SMOTE resampling

Change to XGBoost friendlier format, using all potentially useful columns, dropping Vehicle (an index) and Days (always 365, no chance of separation between loss type events)

```
In [120]: #splitting up the signal and background datasets into training sets, testing sets, and validation sets.
#This splits into 80% to be used for training and 20% testing
X = sim_sum_tot.drop(['Vehicle', 'Days'], axis=1)
y = sim_sum_tot['Loss']
ix = range(y.shape[0])
#X=X.drop('Loss',axis=1)
X_train, X_test, y_train, y_test, ix_train, ix_test = train_test_split(X, y, ix, train_size=0.8)
#X_train, X_val, y_train, y_val, ix_train, ix_val=train_test_split(X_train, y_train, ix_train, test_size=0.2)
```

```
In [121]: X.columns
```

```
Out[121]: Index(['Distance', 'HardBrakes', 'HardAccelerations', 'NightTime_Pct', 'VehicleType', 'Loss'],
              dtype='object')
```

```
In [122]: adasyn = SMOTE(sampling_strategy="auto")
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)
np.bincount(y_train_smote)
```

```
Out[122]: array([20796, 20796], dtype=int64)
```

```
In [123]: X_train_smote
```

```
Out[123]:
```

	Distance	HardBrakes	HardAccelerations	NightTime_Pct	VehicleType	Loss
0	15595	446	16	0.017000	1	1
1	17220	207	155	0.009000	2	0
2	14428	167	2	0.003000	2	0
3	7961	103	1	0.056000	3	0
4	13196	1	10	0.024000	1	0
...	...	...	...	...	...	...
41587	11544	63	142	0.159022	2	1
41588	19066	288	188	0.022176	2	1
41589	14291	77	38	0.078869	2	1
41590	15489	77	386	0.013539	0	1
41591	16749	308	28	0.029469	1	1

41592 rows × 6 columns

```
In [124]: print('Number of training samples: {}'.format(len(X_train)))
print('Number of testing samples: {}'.format(len(X_test)))
print('SMOTE: Number of training samples: {}'.format(len(X_train_smote)))

print('\nNumber of signal events in training set: {}'.format(len(X_train[X_train.Loss == 1])))
print('SMOTE: Number of signal events in training set: {}'.format(len(X_train_smote[X_train_smote.Loss == 1])))
print('Number of background events in training set: {}'.format(len(X_train[X_train.Loss == 0])))
print('Fraction signal: {}'.format(len(X_train[X_train.Loss == 1])/(float)(len(X_train[X_train.Loss == 1]) + len(X_train[X_train.Loss == 0]))))
print('SMOTE: Fraction signal: {}'.format(len(X_train_smote[X_train_smote.Loss == 1])/(float)(len(X_train_adasyn[X_train_smote.Loss == 1]) + len(X_train_smote[X_train_smote.Loss == 0]))))
```

Number of training samples: 24000  
 Number of testing samples: 6000  
 SMOTE: Number of training samples: 41592

Number of signal events in training set: 3204  
 SMOTE: Number of signal events in training set: 20796  
 Number of background events in training set: 20796  
 Fraction signal: 0.1335  
 SMOTE: Fraction signal: 0.5050269561416291

D:\ProgramData\Anaconda3\lib\site-packages\ipykernel\_launcher.py:9: UserWarning: Boolean Series key will be reindexed to match DataFrame index.  
 if \_\_name\_\_ == '\_\_main\_\_':

```
In [125]: features = X_train.columns[:-1] # we skip the last column because it is the Loss Label
features
```

```
Out[125]: Index(['Distance', 'HardBrakes', 'HardAccelerations', 'NightTime_Pct',
                'VehicleType'],
                dtype='object')
```

```
In [126]: binary_bdt_param = {
    "learning_rate" : 0.15,
    "max_depth" : 6,
    "colsample_bytree" : 1.0,
    "subsample" : 1.0,
    "n_estimators" : 200,
    "feature_names" : features,
    'objective' : 'binary:logistic' # objective function
}
binary_task_param = {
    "eval_metric" : ["logloss", "error"],
    "early_stopping_rounds" : 30,
    "eval_set": [(X_train_adasyn[features], y_train_adasyn),
                  (X_test[features], y_test)]
}

binary_bdt = xgb.XGBClassifier(**binary_bdt_param)
binary_bdt.fit(X_train_smote[features], y_train_smote,
               verbose=False, **binary_task_param)
```

[20:07:10] WARNING: C:\Users\Administrator\workspace\xgboost-win64\_release\_1.1.0\src\learner.cc:480:  
Parameters: { feature\_names } might not be used.

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

```
Out[126]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
    colsample_bynode=1, colsample_bytree=1.0,
    feature_names=Index(['Distance', 'HardBrakes', 'HardAcceleratio
ns', 'NightTime_Pct',
    'VehicleType'],
    dtype='object'),
    gamma=0, gpu_id=-1, importance_type='gain',
    interaction_constraints='', learning_rate=0.15, max_delta_step=
0,
    max_depth=6, min_child_weight=1, missing=nan,
    monotone_constraints='()', n_estimators=200, n_jobs=0,
    num_parallel_tree=1, random_state=0, reg_alpha=0, reg_lambda=1,
    scale_pos_weight=1, subsample=1.0, tree_method='exact',
    validate_parameters=1, verbosity=None)
```

```
In [127]: evaluated_df = X_test.copy()
evaluated_df["binary_prob"] = binary_bdt.predict_proba(X_test[features])[:,1]
print(binary_bdt.score(X_test[features],y_test))

binary_bdt.predict_proba(X_test[features])[:,1].round().sum()

0.8045
```

Out[127]: 520.0

```
In [128]: y_pred = binary_bdt.predict(X_test[features])
y_pred.sum()
```

Out[128]: 520

```
In [129]: binary_bdt.predict(X_test[features])
```

Out[129]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)

```
In [130]: from sklearn.metrics import accuracy_score
from sklearn.feature_selection import SelectFromModel


predictions = binary_bdt.predict(X_test[features])
accuracy = accuracy_score(y_test, predictions)
print("Accuracy: %.2f%%" % (accuracy * 100.0))
thresholds = np.sort(binary_bdt.feature_importances_)
for thresh in thresholds:
    # select features using threshold
    selection = SelectFromModel(binary_bdt, threshold=thresh, prefit=True)
    select_X_train = selection.transform(X_train_smote[features])
    # train model
    selection_model = xgb.XGBClassifier()
    selection_model.fit(select_X_train, y_train_smote)
    # eval model
    select_X_test = selection.transform(X_test[features])
    predictions = selection_model.predict(select_X_test)
    accuracy = accuracy_score(y_test, predictions)
    print("Loss Event Accuracy: %.2f%%"%(predictions[y_test==1].sum()/y_test.s
um()*100.))
    print("Thresh=%.3f, n=%d, Accuracy: %.2f%%" % (thresh, select_X_train.shap
e[1], accuracy*100.0))
```

```
Accuracy: 80.45%
Loss Event Accuracy: 12.58%
Thresh=0.060, n=5, Accuracy: 80.30%
Loss Event Accuracy: 10.28%
Thresh=0.062, n=4, Accuracy: 81.22%
Loss Event Accuracy: 10.88%
Thresh=0.084, n=3, Accuracy: 80.97%
Loss Event Accuracy: 4.35%
Thresh=0.273, n=2, Accuracy: 83.18%
Loss Event Accuracy: 3.87%
Thresh=0.521, n=1, Accuracy: 83.95%
```

```
In [131]: predictions[y_test==1].sum()/y_test.sum()
y_pred[y_test==1].sum()/y_test.sum()
```

Out[131]: 0.10519951632406288



```
In [132]: fig, ax_enum = plt.subplots(1,2, figsize=(16,4))
xgb.plot_importance(binary_bdt, importance_type="weight", ax=ax_enum[0], title=
"Weight", show_values=False, grid=False)
xgb.plot_importance(binary_bdt, importance_type="gain", ax=ax_enum[1], title=
"Gain", show_values=False, grid=False)
plt.ylabel("")
plt.sca(ax_enum[1])
plt.ylabel("")
plt.subplots_adjust(wspace=0.3)
```



importance\_type=Gain is how useful is the variable in terms of separation where importance\_type=Weight is how frequently splitting on the variable occurs

```
In [133]: # plot all predictions (both signal and background)
predictions = binary_bdt.predict_proba(X_test[features])[:,1]
plt.figure();
plt.hist(predictions,bins=np.linspace(0,1,50),histtype='step',color='darkgreen',label='All events');
# make the plot readable
plt.xlabel('Prediction from BDT',fontsize=12);
plt.ylabel('Events',fontsize=12);
plt.legend(frameon=False);

# plot signal and background separately
plt.figure();
plt.hist(binary_bdt.predict_proba(X_test[features][y_test==1])[:,1],bins=np.linspace(0,1,50),
         histtype='step',color='midnightblue',label='Loss');
plt.hist(binary_bdt.predict_proba(X_test[features][y_test==0])[:,1],bins=np.linspace(0,1,50),
         histtype='step',color='firebrick',label='No Loss');
# make the plot readable
plt.xlabel('Prediction from BDT',fontsize=12);
plt.ylabel('Events',fontsize=12);
plt.legend(frameon=False);
```



```
In [134]: #Looking at the training set to make sure things still behave, we see they look good
# However the smote events are the ones making up the clearly defined tail
# plot all predictions (both signal and background)
predictions = binary_bdt.predict_proba(X_train_smote[features][:,1])
plt.figure();
plt.hist(predictions,bins=np.linspace(0,1,50),histtype='step',color='darkgreen',label='All events');
# make the plot readable
plt.xlabel('Prediction from BDT',fontsize=12);
plt.ylabel('Events',fontsize=12);
plt.legend(frameon=False);

# plot signal and background separately
plt.figure();
plt.hist(binary_bdt.predict_proba(X_train_smote[features][y_train_smote==1])[:,1],bins=np.linspace(0,1,50),
        histtype='step',color='midnightblue',label='Loss');
plt.hist(binary_bdt.predict_proba(X_train_smote[features][y_train_smote==0])[:,1],bins=np.linspace(0,1,50),
        histtype='step',color='firebrick',label='No Loss');
# make the plot readable
plt.xlabel('Prediction from BDT',fontsize=12);
plt.ylabel('Events',fontsize=12);
plt.legend(frameon=False);
```



